

Introduction

Kaggle Histopathologic Cancer Detection

<https://www.kaggle.com/c/histopathologic-cancer-detection/leaderboard> In this dataset, you are provided with a large number of small pathology images to classify. Files are named with an image id. The train_labels.csv file provides the ground truth for the images in the train folder. You are predicting the labels for the images in the test folder. A positive label indicates that the center 32x32px region of a patch contains at least one pixel of tumor tissue. Tumor tissue in the outer region of the patch does not influence the label. This outer region is provided to enable fully-convolutional models that do not use zero-padding, to ensure consistent behavior when applied to a whole-slide image.

The original PCam dataset contains duplicate images due to its probabilistic sampling, however, the version presented on Kaggle does not contain duplicates. We have otherwise maintained the same data and splits as the PCam benchmark.

The goal of this project is to build a CNN deep learning model that can accurately label images of tumors with cancer. The model will be trained on the training set and evaluated on the test set. The performance of the model will be measured using ROC.

Challenge Problem:

The challenge problem is Histopathologic Cancer Detection, a binary classification task where the goal is to identify metastatic cancer in small image patches taken from larger digital pathology scans. Specifically, the task is to predict whether the center 32x32px region of a patch contains at least one pixel of tumor tissue.

Data Description:

The data consists of a large number of small pathology image patches, each with a size of 96x96 pixels. The images are provided in a train folder and a test folder, with corresponding labels for the train folder provided in a train_labels.csv file. The labels are binary, indicating the presence or absence of tumor tissue in the center 32x32px region of each patch.

- Image Size: 96x96 pixels (RGB)
- Region of Interest: Center 32x32px region
- Labeling: Binary labels indicating presence/absence of tumor tissue in the center region
- Data Structure: Images are stored in a folder with corresponding labels in a CSV file for the train set
- Data Size: 220,025 training images and 57,458 testing images

Due to computation restrictions, the models will be trained and validated on a smaller subsample (1K, 250) of the full training set.

Inputs

```
In [56]: from PIL import Image

import os
import cv2
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import time
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout

from tensorflow.keras.activations import swish
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from sklearn.metrics import roc_auc_score, classification_report, confusion_matrix
from tensorflow.keras.metrics import AUC

In [3]: train_path = './histopathologic-cancer-detection/train'
test_path = './histopathologic-cancer-detection/test'
train_labels = pd.read_csv('./histopathologic-cancer-detection/train_labels.csv')
```


EDA

- Due to computation restriction, a sample of 10K images were selected for EDA.

Label Distribution

```
In [4]: sns.countplot(x='label', data=train_labels)
plt.title('Label Distribution')
plt.xlabel('Label')
plt.ylabel('Count')

Out[4]: Text(0, 0.5, 'Count')
```



Load Data

Reading Image Pixels

- Convert image to RGB
- Normalize by / 255 for 8 bit image
 - 0 (black) to 255 (white)
- Leave as arrays to take advantage of CNN's spatial and data structure awareness
- Append training labels

```
In [5]: # Load file paths and IDs
def load_file_paths(folder_path):
    file_paths = [os.path.join(folder_path, filename) for filename in os.listdir(folder_path)]
    ids = [os.path.basename(file_path).split('.')[0] for file_path in file_paths]
    df = pd.DataFrame({'id': ids, 'file_path': file_paths})
    return df

# Load pixel data on demand
def load_pixel_data(file_path):
    img = cv2.imread(file_path)
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB) # Convert to RGB
    img = img / 255.0 # Normalize pixel values
    return img
```

```
In [ ]: # Load file paths and IDs
train_df = load_file_paths(train_path)
# Join with labels
train_df = pd.merge(train_df, train_labels, on='id')

test_df = load_file_paths(test_path)

sample_df = train_df.sample(n=1000, random_state = 1) # sample train data for testing
```

```
In [7]: # Open an image file
img = Image.open(sample_df['file_path'].iloc[0])

# Get the image size
width, height = img.size

print(f"Image size: {width}x{height}")


Image size: 96x96
```

Overall pixel distribution

```
In [8]: # Load pixel data and plot histogram
train_pixel_values = []
for file_path in sample_df['file_path']:
    img = load_pixel_data(file_path)
    train_pixel_values.extend(img.flatten())

pixel_values = np.array(train_pixel_values)
plt.title('Histogram of Pixel Values in Sampled Images')
plt.hist(pixel_values, bins=50)
plt.xlabel('Pixel Value')
plt.ylabel('Frequency')

Out[8]: Text(0, 0.5, 'Frequency')
```



Brightness:

- The skew left indicates that most pixel values are brighter, with fewer darker pixels.


Contrast:

- The peaks around 1 suggest that the images have a lot of light pixels. (Peak at 1 shows most pixels are close to white)
- The spread of the histogram indicates a moderate contrast, with some variation in pixel brightness.

```
In [9]: # Load pixel data and labels
pixel_values = []
labels = []
for index, row in sample_df.iterrows():
    img = load_pixel_data(row['file_path'])
    pixel_values.extend(img.flatten())
    labels.extend([row['label']] * len(img.flatten()))

# Plot histogram
sns.histplot(x=pixel_values, hue=labels, bins=50)
plt.title('Histogram of Pixel Values in Sampled Images')
plt.xlabel('Pixel Value')
plt.ylabel('Frequency')
```

```
Out[9]: Text(0, 0.5, 'Frequency')
```




The label = 0 distribution shows a continuous increase in pixel concentration as values approach 1, with a huge peak from 0.8 to 1. In contrast, the label = 1 distribution is skewed left with a main peak around 0.7, tapering down towards 0.9 before peaking at 1. Label = 0 has a much higher concentration of pixels near 1 compared to label = 1. The two distributions exhibit distinct patterns, with label = 0 showing a more extreme skew towards bright pixels.

Image Example

```
In [10]: rand_df = train_df.sample(n=9, random_state=1) # Randomly select 9 images
# Create a figure with 3x3 subplots
fig, axes = plt.subplots(3, 3, figsize=(12, 12))

# Load and display each image
for i, file_path in enumerate(rand_df['file_path']):
    img = load_pixel_data(file_path)
    axes[i // 3, i % 3].imshow(img)
    axes[i // 3, i % 3].axis('off')
```




Average Image

```
In [11]: # Load and calculate the average image
avg_img = np.zeros((img.shape[0], img.shape[1], img.shape[2]))
for file_path in sample_df['file_path']:
    img = load_pixel_data(file_path)
    avg_img += img
avg_img /= len(sample_df)

# Display the average image
plt.imshow(avg_img)
```

Out[11]: <matplotlib.image.AxesImage at 0x7fb53f425600>




SD Image

```
In [12]: # Load and calculate the standard deviation image
std_img = np.zeros((img.shape[0], img.shape[1], img.shape[2]))
imgs = []
for file_path in sample_df['file_path']:
    img = load_pixel_data(file_path)
    imgs.append(img)
imgs = np.array(imgs)
std_img = np.std(imgs, axis=0)

# Display the standard deviation image
plt.imshow(std_img)
```

Out[12]: <matplotlib.image.AxesImage at 0x7fb53f4a9480>



EDA Summary

Given the differences in pixel value distributions between the two labels, a Convolutional Neural Network (CNN) model can be a good choice for image classification.

Model Architecture

- Architectures:
 - Simple CNN: A basic CNN with 2-3 convolutional layers, followed by pooling and dense layers. This architecture is a good starting point and can be effective for simple image classification tasks.
 - ResNet CNN: A CNN with residual connections and 5-6 convolutional layers. This architecture can help with vanishing gradients and improve model performance.
- Hyperparameters to Tune:
 - Learning Rate: We'll try different learning rates (e.g., 0.0001, 0.001, 0.01) to see how it affects model convergence and performance.
 - Batch Size: We'll try different batch sizes (e.g., 32, 64, 128) to see how it affects model training time and performance.
 - Number of Convolutional Layers: Future iterations, with more compute, could try different numbers of convolutional layers (e.g., 2, 5) to see how it affects model performance and complexity.
 - Number of Dense Layers: Future iterations, with more compute, could try different numbers of dense layers (e.g., 1, 2) to see how it affects model performance and complexity.
 - Activation Functions: Future iterations, with more compute, could try different activation functions (e.g., ReLU, Leaky ReLU, Swish) to see how it affects model performance.
- Evaluation Metric:
 - ROC-AUC Score: We'll evaluate the performance of different architectures and hyperparameter combinations using the ROC-AUC score, which is the required metric for the Kaggle competition. A higher ROC-AUC score indicates better model performance.

```
In [34]: subsample_df = train_df.sample(n=1250, random_state = 1) # Due to computation
```

```
In [40]: # Define constants
IMG_WIDTH, IMG_HEIGHT = 96, 96
BATCH_SIZE = 64
VALIDATION_SPLIT = 0.2

# Define data generators
train_datagen = ImageDataGenerator(rescale=1./255,
                                   validation_split=VALIDATION_SPLIT)

train_generator = train_datagen.flow_from_dataframe(
    dataframe=subsample_df,
    x_col='file_path',
    y_col='label',
    target_size=(IMG_WIDTH, IMG_HEIGHT),
    batch_size=BATCH_SIZE,
    class_mode='raw',
    subset='training')

validation_generator = train_datagen.flow_from_dataframe(
    dataframe=subsample_df,
    x_col='file_path',
    y_col='label',
    target_size=(IMG_WIDTH, IMG_HEIGHT),
    batch_size=BATCH_SIZE,
    class_mode='raw',
    subset='validation')
```

Found 1000 validated image filenames.
Found 250 validated image filenames.

Simple CNN

Base model

```
In [44]: # Define model architecture
model = Sequential()
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(IMG_WIDTH, IMG_HEIGHT)))
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.2)) # Dropout layer with 20% dropout rate
model.add(Dense(1, activation='sigmoid'))

# Compile model
model.compile(optimizer=Adam(learning_rate=0.001), loss='binary_crossentropy')

# Define early stopping callback
early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=5,
    min_delta=0.001,
    restore_best_weights=True
)

workers = os.cpu_count()
start_time = time.time()

# Train model
history = model.fit(train_generator,
                    epochs=50,
                    validation_data=validation_generator,
                    workers=workers,
                    use_multiprocessing=False,
                    callbacks=[early_stopping])

end_time = time.time()
training_time = end_time - start_time

print(f"Training time: {training_time:.2f} seconds")
```

```

Epoch 1/50
16/16 [=====] - 10s 473ms/step - loss: 0.9906 - auc_
8: 0.5257 - val_loss: 0.6595 - val_auc_8: 0.6369
Epoch 2/50
16/16 [=====] - 8s 483ms/step - loss: 0.6568 - auc_
8: 0.6042 - val_loss: 0.6391 - val_auc_8: 0.8028
Epoch 3/50
16/16 [=====] - 7s 414ms/step - loss: 0.6566 - auc_
8: 0.6042 - val_loss: 0.6294 - val_auc_8: 0.8079
Epoch 4/50
16/16 [=====] - 7s 409ms/step - loss: 0.6316 - auc_
8: 0.6859 - val_loss: 0.6344 - val_auc_8: 0.8484
Epoch 5/50
16/16 [=====] - 7s 417ms/step - loss: 0.6311 - auc_
8: 0.6756 - val_loss: 0.5939 - val_auc_8: 0.8251
Epoch 6/50
16/16 [=====] - 7s 436ms/step - loss: 0.5899 - auc_
8: 0.7945 - val_loss: 0.5563 - val_auc_8: 0.8625
Epoch 7/50
16/16 [=====] - 7s 416ms/step - loss: 0.5558 - auc_
8: 0.8205 - val_loss: 0.5577 - val_auc_8: 0.8327
Epoch 8/50
16/16 [=====] - 7s 410ms/step - loss: 0.5089 - auc_
8: 0.8331 - val_loss: 0.4727 - val_auc_8: 0.8588
Epoch 9/50
16/16 [=====] - 7s 409ms/step - loss: 0.4927 - auc_
8: 0.8354 - val_loss: 0.4686 - val_auc_8: 0.8625
Epoch 10/50
16/16 [=====] - 7s 413ms/step - loss: 0.4357 - auc_
8: 0.8772 - val_loss: 0.4723 - val_auc_8: 0.8515
Epoch 11/50
16/16 [=====] - 7s 428ms/step - loss: 0.4230 - auc_
8: 0.8901 - val_loss: 0.5317 - val_auc_8: 0.8261
Epoch 12/50
16/16 [=====] - 7s 407ms/step - loss: 0.4028 - auc_
8: 0.9018 - val_loss: 0.5040 - val_auc_8: 0.8574
Epoch 13/50
16/16 [=====] - 7s 412ms/step - loss: 0.3349 - auc_
8: 0.9392 - val_loss: 0.4735 - val_auc_8: 0.8451
Epoch 14/50
16/16 [=====] - 7s 414ms/step - loss: 0.2927 - auc_
8: 0.9497 - val_loss: 0.5447 - val_auc_8: 0.8295
Training time: 101.13 seconds

```

```

In [ ]: # Plot loss
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend()
plt.show()


# Calculate ROC-AUC score
y_pred_train = model.predict(train_generator)
y_pred_val = model.predict(validation_generator)
y_train = train_generator.labels
y_val = validation_generator.labels

train_roc_auc = roc_auc_score(y_train, y_pred_train.flatten())
val_roc_auc = roc_auc_score(y_val, y_pred_val.flatten())

print(f'Training ROC-AUC: {train_roc_auc:.4f}')
print(f'Validation ROC-AUC: {val_roc_auc:.4f}')

# Calculate classification matrix
y_pred_val_class = (y_pred_val > 0.5).astype('int32') # default threshold of 0.5
print(ClassificationReport(y_val, y_pred_val_class))
print(confusion_matrix(y_val, y_pred_val_class))

```



```

16/16 [=====] - 2s 133ms/step
4/4 [=====] - 1s 112ms/step
Training ROC-AUC: 0.5088
Validation ROC-AUC: 0.4998

```

	precision	recall	f1-score	support
0	0.61	0.70	0.65	151
1	0.42	0.33	0.37	99
accuracy			0.55	250
macro avg	0.52	0.51	0.51	250
weighted avg	0.54	0.55	0.54	250

```

[[105 46]
 [ 66 33]]

```

The baseline CNN model achieved a validation ROC-AUC score of 0.4998, which is essentially at chance level, but the model's performance metrics suggest it still has some ability to distinguish between classes. However, the learning curve shows signs of overfitting, with divergence between training and validation loss starting around epoch 10, indicating that the model may not generalize well to new data. Despite this, the model's performance is still relatively good, with an accuracy of 0.55 and decent precision and recall for class 0.

Hyper parameter tuning

- Batch Size
- Learning Rate

```

In [49]: # Define hyperparameter grid
batch_sizes = [32, 64, 128]
learning_rates = [0.0001, 0.001, 0.01]

# Create a list to store results
results = []

for batch_size in batch_sizes:
    for learning_rate in learning_rates:
        # Define data generators
        train_generator = train_datagen.flow_from_dataframe(
            dataframe=subsample_df,
            x_col='file_path',
            y_col='label',
            target_size=(IMG_WIDTH, IMG_HEIGHT),
            batch_size=batch_size,
            class_mode='raw',
            subset='training')

        validation_generator = train_datagen.flow_from_dataframe(
            dataframe=subsample_df,
            x_col='file_path',
            y_col='label',
            target_size=(IMG_WIDTH, IMG_HEIGHT),
            batch_size=batch_size,
            class_mode='raw',
            subset='validation')

        # Define model architecture
        model = Sequential()
        model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(IMG_WIDTH, IMG_HEIGHT, 3)))
        model.add(MaxPooling2D((2, 2)))
        model.add(Conv2D(64, (3, 3), activation='relu'))
        model.add(MaxPooling2D((2, 2)))
        model.add(Flatten())
        model.add(Dense(128, activation='relu'))
        model.add(Dropout(0.2))
        model.add(Dense(1, activation='sigmoid'))

        # Compile model
        model.compile(optimizer=Adam(learning_rate=learning_rate), loss='binary_crossentropy')

        # Define early stopping callback
        early_stopping = EarlyStopping(
            monitor='val_loss',
            patience=5,
            min_delta=0.001,
            restore_best_weights=True)

        # Train model
        start_time = time.time()
        history = model.fit(train_generator,
                            epochs=50,
                            validation_data=validation_generator,
                            workers=workers,
                            use_multiprocessing=False,
                            callbacks=[early_stopping])
        end_time = time.time()
        training_time = end_time - start_time

        # Calculate ROC-AUC score
        y_pred_train = model.predict(train_generator)
        y_pred_val = model.predict(validation_generator)
        y_train = train_generator.labels
        y_val = validation_generator.labels

        train_roc_auc = roc_auc_score(y_train, y_pred_train.flatten())
        val_roc_auc = roc_auc_score(y_val, y_pred_val.flatten())

        # Store results
        row_name = f'batch{batch_size}_lr{str(learning_rate).replace(".", "")}'
        results.append({
            'model_name': row_name,
            'batch_size': batch_size,
            'learning_rate': learning_rate,
            'train_time': training_time,
            'train_roc_auc': train_roc_auc,
            'val_roc_auc': val_roc_auc
        })

```



```
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
32/32 [=====] - 9s 253ms/step - loss: 0.6723 - auc_
18: 0.5596 - val_loss: 0.6252 - val_auc_18: 0.8249
Epoch 2/50
32/32 [=====] - 8s 236ms/step - loss: 0.6214 - auc_
18: 0.7112 - val_loss: 0.5797 - val_auc_18: 0.8407
Epoch 3/50
32/32 [=====] - 13s 393ms/step - loss: 0.5823 - auc_
18: 0.7792 - val_loss: 0.5391 - val_auc_18: 0.8545
Epoch 4/50
32/32 [=====] - 7s 222ms/step - loss: 0.5517 - auc_
18: 0.7940 - val_loss: 0.5149 - val_auc_18: 0.8461
Epoch 5/50
32/32 [=====] - 8s 235ms/step - loss: 0.5282 - auc_
18: 0.8153 - val_loss: 0.5344 - val_auc_18: 0.8607
Epoch 6/50
32/32 [=====] - 8s 247ms/step - loss: 0.5115 - auc_
18: 0.8239 - val_loss: 0.5096 - val_auc_18: 0.8568
Epoch 7/50
32/32 [=====] - 8s 254ms/step - loss: 0.4930 - auc_
18: 0.8422 - val_loss: 0.4650 - val_auc_18: 0.8648
Epoch 8/50
32/32 [=====] - 8s 246ms/step - loss: 0.4798 - auc_
18: 0.8501 - val_loss: 0.4771 - val_auc_18: 0.8652
Epoch 9/50
32/32 [=====] - 8s 243ms/step - loss: 0.4577 - auc_
18: 0.8674 - val_loss: 0.4506 - val_auc_18: 0.8687
Epoch 10/50
32/32 [=====] - 8s 246ms/step - loss: 0.4623 - auc_
18: 0.8589 - val_loss: 0.4499 - val_auc_18: 0.8676
Epoch 11/50
32/32 [=====] - 8s 242ms/step - loss: 0.4467 - auc_
18: 0.8711 - val_loss: 0.4911 - val_auc_18: 0.8658
Epoch 12/50
32/32 [=====] - 8s 246ms/step - loss: 0.4282 - auc_
18: 0.8850 - val_loss: 0.4458 - val_auc_18: 0.8690
Epoch 13/50
32/32 [=====] - 8s 241ms/step - loss: 0.4115 - auc_
18: 0.8985 - val_loss: 0.4443 - val_auc_18: 0.8694
Epoch 14/50
32/32 [=====] - 8s 242ms/step - loss: 0.3912 - auc_
18: 0.9110 - val_loss: 0.4579 - val_auc_18: 0.8678
Epoch 15/50
32/32 [=====] - 10s 311ms/step - loss: 0.3765 - auc_
18: 0.9211 - val_loss: 0.4558 - val_auc_18: 0.8711
Epoch 16/50
32/32 [=====] - 8s 253ms/step - loss: 0.3622 - auc_
18: 0.9268 - val_loss: 0.4404 - val_auc_18: 0.8699
Epoch 17/50
32/32 [=====] - 8s 247ms/step - loss: 0.3387 - auc_
18: 0.9402 - val_loss: 0.4403 - val_auc_18: 0.8693
Epoch 18/50
32/32 [=====] - 8s 258ms/step - loss: 0.3387 - auc_
18: 0.9363 - val_loss: 0.4545 - val_auc_18: 0.8602
Epoch 19/50
32/32 [=====] - 8s 246ms/step - loss: 0.3292 - auc_
18: 0.9404 - val_loss: 0.4489 - val_auc_18: 0.8712
Epoch 20/50
32/32 [=====] - 8s 257ms/step - loss: 0.3051 - auc_
18: 0.9551 - val_loss: 0.4634 - val_auc_18: 0.8689
Epoch 21/50
32/32 [=====] - 8s 246ms/step - loss: 0.2932 - auc_
18: 0.9580 - val_loss: 0.4567 - val_auc_18: 0.8677
32/32 [=====] - 2s 73ms/step
8/8 [=====] - 1s 72ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
32/32 [=====] - 10s 257ms/step - loss: 0.8164 - auc_
19: 0.5662 - val_loss: 0.5822 - val_auc_19: 0.8203
Epoch 2/50
32/32 [=====] - 8s 247ms/step - loss: 0.5729 - auc_
19: 0.7693 - val_loss: 0.6120 - val_auc_19: 0.8533
Epoch 3/50
32/32 [=====] - 8s 255ms/step - loss: 0.5453 - auc_
19: 0.7845 - val_loss: 0.6174 - val_auc_19: 0.7834
Epoch 4/50
32/32 [=====] - 8s 251ms/step - loss: 0.5187 - auc_
19: 0.8123 - val_loss: 0.4573 - val_auc_19: 0.8702
Epoch 5/50
32/32 [=====] - 8s 257ms/step - loss: 0.4696 - auc_
19: 0.8535 - val_loss: 0.4905 - val_auc_19: 0.8688
Epoch 6/50
32/32 [=====] - 9s 264ms/step - loss: 0.4351 - auc_
19: 0.8797 - val_loss: 0.4682 - val_auc_19: 0.8515
Epoch 7/50
32/32 [=====] - 8s 255ms/step - loss: 0.3473 - auc_
19: 0.9265 - val_loss: 0.5793 - val_auc_19: 0.8543
Epoch 8/50
32/32 [=====] - 14s 422ms/step - loss: 0.3121 - auc_
19: 0.9406 - val_loss: 0.5165 - val_auc_19: 0.8308
Epoch 9/50
32/32 [=====] - 12s 390ms/step - loss: 0.2412 - auc_
19: 0.9648 - val_loss: 0.5049 - val_auc_19: 0.8396
32/32 [=====] - 2s 68ms/step
8/8 [=====] - 1s 66ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
32/32 [=====] - 9s 250ms/step - loss: 3.3487 - auc_
20: 0.5183 - val_loss: 0.6318 - val_auc_20: 0.7544
Epoch 2/50
32/32 [=====] - 8s 242ms/step - loss: 0.6529 - auc_
20: 0.5895 - val_loss: 0.6824 - val_auc_20: 0.7458
Epoch 3/50
32/32 [=====] - 8s 252ms/step - loss: 0.6395 - auc_
20: 0.6347 - val_loss: 0.6488 - val_auc_20: 0.7927
Epoch 4/50
32/32 [=====] - 8s 248ms/step - loss: 0.6355 - auc_
20: 0.6651 - val_loss: 0.6109 - val_auc_20: 0.6975
Epoch 5/50
32/32 [=====] - 8s 247ms/step - loss: 0.6545 - auc_
```

```
20: 0.6383 - val_loss: 0.6269 - val_auc_20: 0.6954
Epoch 6/50
32/32 [=====] - 8s 245ms/step - loss: 0.6323 - auc_
20: 0.6389 - val_loss: 0.6967 - val_auc_20: 0.5531
Epoch 7/50
32/32 [=====] - 8s 248ms/step - loss: 0.6780 - auc_
20: 0.4660 - val_loss: 0.6706 - val_auc_20: 0.6632
Epoch 8/50
32/32 [=====] - 8s 253ms/step - loss: 0.6637 - auc_
20: 0.5673 - val_loss: 0.6854 - val_auc_20: 0.4807
Epoch 9/50
32/32 [=====] - 8s 242ms/step - loss: 0.6622 - auc_
20: 0.5771 - val_loss: 0.6343 - val_auc_20: 0.6711
32/32 [=====] - 3s 82ms/step
8/8 [=====] - 1s 69ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
16/16 [=====] - 9s 492ms/step - loss: 0.6985 - auc_
21: 0.5200 - val_loss: 0.6666 - val_auc_21: 0.7240
Epoch 2/50
16/16 [=====] - 8s 484ms/step - loss: 0.6559 - auc_
21: 0.6206 - val_loss: 0.6432 - val_auc_21: 0.8019
Epoch 3/50
16/16 [=====] - 8s 479ms/step - loss: 0.6435 - auc_
21: 0.6638 - val_loss: 0.6170 - val_auc_21: 0.8383
Epoch 4/50
16/16 [=====] - 8s 481ms/step - loss: 0.6224 - auc_
21: 0.7194 - val_loss: 0.5986 - val_auc_21: 0.8489
Epoch 5/50
16/16 [=====] - 8s 484ms/step - loss: 0.6006 - auc_
21: 0.7630 - val_loss: 0.5783 - val_auc_21: 0.8563
Epoch 6/50
16/16 [=====] - 8s 498ms/step - loss: 0.5802 - auc_
21: 0.7905 - val_loss: 0.5489 - val_auc_21: 0.8611
Epoch 7/50
16/16 [=====] - 8s 484ms/step - loss: 0.5561 - auc_
21: 0.8181 - val_loss: 0.5242 - val_auc_21: 0.8634
Epoch 8/50
16/16 [=====] - 8s 491ms/step - loss: 0.5486 - auc_
21: 0.8015 - val_loss: 0.5025 - val_auc_21: 0.8635
Epoch 9/50
16/16 [=====] - 8s 501ms/step - loss: 0.5282 - auc_
21: 0.8203 - val_loss: 0.4910 - val_auc_21: 0.8635
Epoch 10/50
16/16 [=====] - 9s 504ms/step - loss: 0.5158 - auc_
21: 0.8283 - val_loss: 0.4776 - val_auc_21: 0.8641
Epoch 11/50
16/16 [=====] - 8s 476ms/step - loss: 0.5068 - auc_
21: 0.8295 - val_loss: 0.4750 - val_auc_21: 0.8654
Epoch 12/50
16/16 [=====] - 8s 493ms/step - loss: 0.5025 - auc_
21: 0.8306 - val_loss: 0.4634 - val_auc_21: 0.8640
Epoch 13/50
16/16 [=====] - 8s 469ms/step - loss: 0.4975 - auc_
21: 0.8322 - val_loss: 0.4991 - val_auc_21: 0.8638
Epoch 14/50
16/16 [=====] - 8s 493ms/step - loss: 0.5241 - auc_
21: 0.8083 - val_loss: 0.4877 - val_auc_21: 0.8637
Epoch 15/50
16/16 [=====] - 8s 476ms/step - loss: 0.5082 - auc_
21: 0.8227 - val_loss: 0.4651 - val_auc_21: 0.8632
Epoch 16/50
16/16 [=====] - 8s 481ms/step - loss: 0.4924 - auc_
21: 0.8362 - val_loss: 0.4513 - val_auc_21: 0.8656
Epoch 17/50
16/16 [=====] - 8s 479ms/step - loss: 0.4787 - auc_
21: 0.8463 - val_loss: 0.4555 - val_auc_21: 0.8675
Epoch 18/50
16/16 [=====] - 8s 497ms/step - loss: 0.4705 - auc_
21: 0.8532 - val_loss: 0.4486 - val_auc_21: 0.8639
Epoch 19/50
16/16 [=====] - 8s 474ms/step - loss: 0.4675 - auc_
21: 0.8544 - val_loss: 0.4564 - val_auc_21: 0.8663
Epoch 20/50
16/16 [=====] - 8s 469ms/step - loss: 0.4665 - auc_
21: 0.8540 - val_loss: 0.4477 - val_auc_21: 0.8645
Epoch 21/50
16/16 [=====] - 8s 477ms/step - loss: 0.4572 - auc_
21: 0.8619 - val_loss: 0.4528 - val_auc_21: 0.8656
Epoch 22/50
16/16 [=====] - 8s 491ms/step - loss: 0.4627 - auc_
21: 0.8575 - val_loss: 0.4585 - val_auc_21: 0.8685
Epoch 23/50
16/16 [=====] - 8s 476ms/step - loss: 0.4669 - auc_
21: 0.8532 - val_loss: 0.4524 - val_auc_21: 0.8668
16/16 [=====] - 2s 125ms/step
4/4 [=====] - 1s 126ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
16/16 [=====] - 9s 498ms/step - loss: 1.1570 - auc_
22: 0.5531 - val_loss: 0.6398 - val_auc_22: 0.7650
Epoch 2/50
16/16 [=====] - 8s 470ms/step - loss: 0.6261 - auc_
22: 0.6927 - val_loss: 0.5957 - val_auc_22: 0.8278
Epoch 3/50
16/16 [=====] - 8s 486ms/step - loss: 0.5994 - auc_
22: 0.7681 - val_loss: 0.5924 - val_auc_22: 0.7730
Epoch 4/50
16/16 [=====] - 8s 466ms/step - loss: 0.5768 - auc_
22: 0.7794 - val_loss: 0.5268 - val_auc_22: 0.8654
Epoch 5/50
16/16 [=====] - 8s 471ms/step - loss: 0.5382 - auc_
22: 0.7958 - val_loss: 0.4779 - val_auc_22: 0.8589
Epoch 6/50
16/16 [=====] - 8s 490ms/step - loss: 0.4933 - auc_
22: 0.8428 - val_loss: 0.4578 - val_auc_22: 0.8603
Epoch 7/50
16/16 [=====] - 9s 514ms/step - loss: 0.4446 - auc_
22: 0.8662 - val_loss: 0.4833 - val_auc_22: 0.8621
Epoch 8/50
16/16 [=====] - 8s 474ms/step - loss: 0.4223 - auc_
22: 0.8857 - val_loss: 0.4708 - val_auc_22: 0.8547
```

```
Epoch 9/50
16/16 [=====] - 7s 445ms/step - loss: 0.4038 - auc_
22: 0.9010 - val_loss: 0.5046 - val_auc_22: 0.8445
Epoch 10/50
16/16 [=====] - 8s 477ms/step - loss: 0.3680 - auc_
22: 0.9158 - val_loss: 0.4896 - val_auc_22: 0.8458
Epoch 11/50
16/16 [=====] - 8s 477ms/step - loss: 0.3250 - auc_
22: 0.9398 - val_loss: 0.4826 - val_auc_22: 0.8468
16/16 [=====] - 2s 133ms/step
4/4 [=====] - 1s 120ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
16/16 [=====] - 10s 526ms/step - loss: 7.0240 - auc_
23: 0.5223 - val_loss: 0.6898 - val_auc_23: 0.6110
Epoch 2/50
16/16 [=====] - 8s 469ms/step - loss: 0.6855 - auc_
23: 0.5218 - val_loss: 0.6695 - val_auc_23: 0.5739
Epoch 3/50
16/16 [=====] - 8s 462ms/step - loss: 0.6733 - auc_
23: 0.5091 - val_loss: 0.6697 - val_auc_23: 0.5761
Epoch 4/50
16/16 [=====] - 8s 476ms/step - loss: 0.6737 - auc_
23: 0.4973 - val_loss: 0.6694 - val_auc_23: 0.5696
Epoch 5/50
16/16 [=====] - 8s 469ms/step - loss: 0.6738 - auc_
23: 0.5082 - val_loss: 0.6738 - val_auc_23: 0.5532
Epoch 6/50
16/16 [=====] - 8s 484ms/step - loss: 0.6699 - auc_
23: 0.5368 - val_loss: 0.6736 - val_auc_23: 0.5804
Epoch 7/50
16/16 [=====] - 8s 485ms/step - loss: 0.6792 - auc_
23: 0.4711 - val_loss: 0.6726 - val_auc_23: 0.5949
16/16 [=====] - 2s 122ms/step
4/4 [=====] - 1s 127ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
8/8 [=====] - 11s 1s/step - loss: 0.7366 - auc_24:
0.4991 - val_loss: 0.6694 - val_auc_24: 0.6724
Epoch 2/50
8/8 [=====] - 9s 1s/step - loss: 0.6663 - auc_24:
0.5663 - val_loss: 0.6602 - val_auc_24: 0.7478
Epoch 3/50
8/8 [=====] - 8s 938ms/step - loss: 0.6635 - auc_2
4: 0.5727 - val_loss: 0.6493 - val_auc_24: 0.7813
Epoch 4/50
8/8 [=====] - 8s 942ms/step - loss: 0.6518 - auc_2
4: 0.6510 - val_loss: 0.6383 - val_auc_24: 0.7814
Epoch 5/50
8/8 [=====] - 9s 999ms/step - loss: 0.6415 - auc_2
4: 0.6887 - val_loss: 0.6266 - val_auc_24: 0.7886
Epoch 6/50
8/8 [=====] - 9s 997ms/step - loss: 0.6309 - auc_2
4: 0.7234 - val_loss: 0.6124 - val_auc_24: 0.8033
Epoch 7/50
8/8 [=====] - 9s 1s/step - loss: 0.6191 - auc_24:
0.7460 - val_loss: 0.5987 - val_auc_24: 0.8188
Epoch 8/50
8/8 [=====] - 8s 990ms/step - loss: 0.6077 - auc_2
4: 0.7703 - val_loss: 0.5846 - val_auc_24: 0.8336
Epoch 9/50
8/8 [=====] - 8s 990ms/step - loss: 0.5961 - auc_2
4: 0.7700 - val_loss: 0.5814 - val_auc_24: 0.8377
Epoch 10/50
8/8 [=====] - 9s 1s/step - loss: 0.5829 - auc_24:
0.7858 - val_loss: 0.5589 - val_auc_24: 0.8494
Epoch 11/50
8/8 [=====] - 9s 1s/step - loss: 0.5711 - auc_24:
0.7986 - val_loss: 0.5453 - val_auc_24: 0.8555
Epoch 12/50
8/8 [=====] - 9s 1s/step - loss: 0.5576 - auc_24:
0.8161 - val_loss: 0.5289 - val_auc_24: 0.8581
Epoch 13/50
8/8 [=====] - 9s 1s/step - loss: 0.5428 - auc_24:
0.8247 - val_loss: 0.5170 - val_auc_24: 0.8607
Epoch 14/50
8/8 [=====] - 9s 997ms/step - loss: 0.5346 - auc_2
4: 0.8307 - val_loss: 0.5075 - val_auc_24: 0.8627
Epoch 15/50
8/8 [=====] - 8s 947ms/step - loss: 0.5215 - auc_2
4: 0.8390 - val_loss: 0.5055 - val_auc_24: 0.8633
Epoch 16/50
8/8 [=====] - 9s 1s/step - loss: 0.5176 - auc_24:
0.8320 - val_loss: 0.4987 - val_auc_24: 0.8621
Epoch 17/50
8/8 [=====] - 9s 987ms/step - loss: 0.5122 - auc_2
4: 0.8316 - val_loss: 0.4904 - val_auc_24: 0.8632
Epoch 18/50
8/8 [=====] - 9s 1s/step - loss: 0.4951 - auc_24:
0.8432 - val_loss: 0.4739 - val_auc_24: 0.8637
Epoch 19/50
8/8 [=====] - 8s 971ms/step - loss: 0.4910 - auc_2
4: 0.8490 - val_loss: 0.4739 - val_auc_24: 0.8627
Epoch 20/50
8/8 [=====] - 9s 1s/step - loss: 0.4789 - auc_24:
0.8552 - val_loss: 0.4669 - val_auc_24: 0.8632
Epoch 21/50
8/8 [=====] - 9s 1s/step - loss: 0.4855 - auc_24:
0.8480 - val_loss: 0.4643 - val_auc_24: 0.8631
Epoch 22/50
8/8 [=====] - 8s 1s/step - loss: 0.4915 - auc_24:
0.8384 - val_loss: 0.4737 - val_auc_24: 0.8638
Epoch 23/50
8/8 [=====] - 8s 1s/step - loss: 0.4744 - auc_24:
0.8547 - val_loss: 0.4818 - val_auc_24: 0.8625
Epoch 24/50
8/8 [=====] - 8s 1s/step - loss: 0.4653 - auc_24:
0.8616 - val_loss: 0.4612 - val_auc_24: 0.8619
Epoch 25/50
8/8 [=====] - 9s 1s/step - loss: 0.4674 - auc_24:
0.8586 - val_loss: 0.4600 - val_auc_24: 0.8622
Epoch 26/50
```

```

8/8 [=====] - 9s 1s/step - loss: 0.4658 - auc_24:
0.8584 - val_loss: 0.4577 - val_auc_24: 0.8628
Epoch 27/50
8/8 [=====] - 9s 1s/step - loss: 0.4574 - auc_24:
0.8663 - val_loss: 0.4573 - val_auc_24: 0.8624
Epoch 28/50
8/8 [=====] - 9s 1s/step - loss: 0.4471 - auc_24:
0.8754 - val_loss: 0.4681 - val_auc_24: 0.8638
Epoch 29/50
8/8 [=====] - 9s 1s/step - loss: 0.4485 - auc_24:
0.8705 - val_loss: 0.4620 - val_auc_24: 0.8629
Epoch 30/50
8/8 [=====] - 9s 1s/step - loss: 0.4433 - auc_24:
0.8754 - val_loss: 0.4532 - val_auc_24: 0.8629
Epoch 31/50
8/8 [=====] - 8s 980ms/step - loss: 0.4391 - auc_2
4: 0.8793 - val_loss: 0.4585 - val_auc_24: 0.8633
Epoch 32/50
8/8 [=====] - 9s 1s/step - loss: 0.4354 - auc_24:
0.8821 - val_loss: 0.4760 - val_auc_24: 0.8579
Epoch 33/50
8/8 [=====] - 9s 1s/step - loss: 0.4401 - auc_24:
0.8759 - val_loss: 0.4583 - val_auc_24: 0.8615
Epoch 34/50
8/8 [=====] - 10s 1s/step - loss: 0.4317 - auc_24:
0.8847 - val_loss: 0.4812 - val_auc_24: 0.8630
Epoch 35/50
8/8 [=====] - 9s 1s/step - loss: 0.4248 - auc_24:
0.8881 - val_loss: 0.4631 - val_auc_24: 0.8621
8/8 [=====] - 2s 250ms/step
2/2 [=====] - 1s 243ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
8/8 [=====] - 13s 1s/step - loss: 1.2164 - auc_25:
0.5221 - val_loss: 0.6635 - val_auc_25: 0.6837
Epoch 2/50
8/8 [=====] - 8s 979ms/step - loss: 0.6594 - auc_2
5: 0.6022 - val_loss: 0.6256 - val_auc_25: 0.8276
Epoch 3/50
8/8 [=====] - 9s 1s/step - loss: 0.6264 - auc_25:
0.7332 - val_loss: 0.5859 - val_auc_25: 0.8596
Epoch 4/50
8/8 [=====] - 8s 911ms/step - loss: 0.5900 - auc_2
5: 0.8145 - val_loss: 0.5668 - val_auc_25: 0.8589
Epoch 5/50
8/8 [=====] - 9s 995ms/step - loss: 0.5701 - auc_2
5: 0.8295 - val_loss: 0.5471 - val_auc_25: 0.8657
Epoch 6/50
8/8 [=====] - 9s 971ms/step - loss: 0.5447 - auc_2
5: 0.8524 - val_loss: 0.5340 - val_auc_25: 0.8664
Epoch 7/50
8/8 [=====] - 8s 974ms/step - loss: 0.5407 - auc_2
5: 0.8405 - val_loss: 0.5434 - val_auc_25: 0.8415
Epoch 8/50
8/8 [=====] - 8s 968ms/step - loss: 0.4942 - auc_2
5: 0.8450 - val_loss: 0.4546 - val_auc_25: 0.8647
Epoch 9/50
8/8 [=====] - 8s 987ms/step - loss: 0.4607 - auc_2
5: 0.8592 - val_loss: 0.4714 - val_auc_25: 0.8543
Epoch 10/50
8/8 [=====] - 8s 980ms/step - loss: 0.3986 - auc_2
5: 0.9048 - val_loss: 0.4621 - val_auc_25: 0.8562
Epoch 11/50
8/8 [=====] - 8s 909ms/step - loss: 0.3671 - auc_2
5: 0.9192 - val_loss: 0.4648 - val_auc_25: 0.8552
Epoch 12/50
8/8 [=====] - 8s 990ms/step - loss: 0.3316 - auc_2
5: 0.9357 - val_loss: 0.4808 - val_auc_25: 0.8497
Epoch 13/50
8/8 [=====] - 8s 927ms/step - loss: 0.2899 - auc_2
5: 0.9522 - val_loss: 0.4963 - val_auc_25: 0.8418
8/8 [=====] - 2s 244ms/step
2/2 [=====] - 1s 238ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
8/8 [=====] - 10s 1s/step - loss: 8.7792 - auc_26:
0.5141 - val_loss: 0.6993 - val_auc_26: 0.5157
Epoch 2/50
8/8 [=====] - 9s 992ms/step - loss: 0.6860 - auc_2
6: 0.4932 - val_loss: 0.6803 - val_auc_26: 0.5000
Epoch 3/50
8/8 [=====] - 8s 978ms/step - loss: 0.6774 - auc_2
6: 0.5032 - val_loss: 0.6750 - val_auc_26: 0.5000
Epoch 4/50
8/8 [=====] - 8s 960ms/step - loss: 0.6714 - auc_2
6: 0.5279 - val_loss: 0.6719 - val_auc_26: 0.5000
Epoch 5/50
8/8 [=====] - 8s 900ms/step - loss: 0.6706 - auc_2
6: 0.4966 - val_loss: 0.6715 - val_auc_26: 0.5000
Epoch 6/50
8/8 [=====] - 8s 909ms/step - loss: 0.6708 - auc_2
6: 0.4933 - val_loss: 0.6717 - val_auc_26: 0.5000
Epoch 7/50
8/8 [=====] - 8s 886ms/step - loss: 0.6706 - auc_2
6: 0.4987 - val_loss: 0.6716 - val_auc_26: 0.5000
Epoch 8/50
8/8 [=====] - 9s 1s/step - loss: 0.6702 - auc_26:
0.5007 - val_loss: 0.6714 - val_auc_26: 0.5000
Epoch 9/50
8/8 [=====] - 8s 894ms/step - loss: 0.6710 - auc_2
6: 0.4763 - val_loss: 0.6714 - val_auc_26: 0.5000
8/8 [=====] - 2s 249ms/step
2/2 [=====] - 1s 233ms/step

```

```

In [50]: cnn_results = pd.DataFrame(results).set_index('model_name')
cnn_results

```

Out[50]:

	batch_size	learning_rate	train_time	train_roc_auc	val_roc_auc
--	------------	---------------	------------	---------------	-------------

model_name					
batch32_lr00001	32	0.0001	179.060640	0.485649	0.463041
batch32_lr0001	32	0.0010	85.950364	0.498820	0.491939
batch32_lr001	32	0.0100	74.154012	0.471669	0.489431
batch64_lr00001	64	0.0001	188.133463	0.514129	0.499699
batch64_lr0001	64	0.0010	91.393906	0.501683	0.461770
batch64_lr001	64	0.0100	57.437663	0.499067	0.503378
batch128_lr00001	128	0.0001	313.107621	0.517441	0.446117
batch128_lr0001	128	0.0010	113.998086	0.488529	0.529534
batch128_lr001	128	0.0100	76.276346	0.500000	0.500000

The results show that the best validation ROC-AUC score is achieved with a batch size of 128 and a learning rate of 0.001, with a score of 0.529534. The training time generally decreases as the learning rate increases, suggesting that higher learning rates lead to faster convergence. However, the training time also varies with batch size, with batch size 64 resulting in faster training times than batch size 32 or 128. The validation ROC-AUC scores are generally close to 0.5, indicating that the models are not strongly distinguishing between classes. Overall, the results suggest that a batch size of 128 and a learning rate of 0.001 may be a good combination for this model, but further tuning may be needed to achieve better performance.

```
In [65]: # Define constants
IMG_WIDTH, IMG_HEIGHT = 96, 96
BATCH_SIZE = 128
VALIDATION_SPLIT = 0.2

# Define data generators
train_datagen = ImageDataGenerator(rescale=1./255,
                                   validation_split=VALIDATION_SPLIT)

train_generator = train_datagen.flow_from_dataframe(
    dataframe=subsample_df,
    x_col='file_path',
    y_col='label',
    target_size=(IMG_WIDTH, IMG_HEIGHT),
    batch_size=BATCH_SIZE,
    class_mode='raw',
    subset='training')

validation_generator = train_datagen.flow_from_dataframe(
    dataframe=subsample_df,
    x_col='file_path',
    y_col='label',
    target_size=(IMG_WIDTH, IMG_HEIGHT),
    batch_size=BATCH_SIZE,
    class_mode='raw',
    subset='validation')

# Define cnn_final architecture
cnn_final = Sequential()
cnn_final.add(Conv2D(32, (3, 3), activation='relu', input_shape=(IMG_WIDTH,
cnn_final.add(MaxPooling2D((2, 2)))
cnn_final.add(Conv2D(64, (3, 3), activation='relu'))
cnn_final.add(MaxPooling2D((2, 2)))
cnn_final.add(Flatten())
cnn_final.add(Dense(128, activation='relu'))
cnn_final.add(Dropout(0.2)) # Dropout layer with 20% dropout rate
cnn_final.add(Dense(1, activation='sigmoid'))

# Compile cnn_final
cnn_final.compile(optimizer=Adam(learning_rate=0.001), loss='binary_crossentropy')

start_time = time.time()
# Train cnn_final
history = cnn_final.fit(train_generator,
                        epochs=50,
                        validation_data=validation_generator,
                        workers=workers,
                        use_multiprocessing=False,
                        callbacks=[early_stopping])
end_time = time.time()
training_time = end_time - start_time

print(f"Training time: {training_time:.2f} seconds")
```

```

Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
8/8 [=====] - 9s 873ms/step - loss: 1.4339 - auc_3
9: 0.4766 - val_loss: 0.7184 - val_auc_39: 0.7369
Epoch 2/50
8/8 [=====] - 7s 802ms/step - loss: 0.6748 - auc_3
9: 0.5517 - val_loss: 0.6372 - val_auc_39: 0.7048
Epoch 3/50
8/8 [=====] - 7s 849ms/step - loss: 0.6287 - auc_3
9: 0.7145 - val_loss: 0.5924 - val_auc_39: 0.8283
Epoch 4/50
8/8 [=====] - 8s 888ms/step - loss: 0.6016 - auc_3
9: 0.7720 - val_loss: 0.5684 - val_auc_39: 0.8462
Epoch 5/50
8/8 [=====] - 8s 933ms/step - loss: 0.5954 - auc_3
9: 0.7699 - val_loss: 0.6037 - val_auc_39: 0.8430
Epoch 6/50
8/8 [=====] - 7s 865ms/step - loss: 0.5720 - auc_3
9: 0.8030 - val_loss: 0.5567 - val_auc_39: 0.8725
Epoch 7/50
8/8 [=====] - 7s 877ms/step - loss: 0.5377 - auc_3
9: 0.8517 - val_loss: 0.5321 - val_auc_39: 0.8632
Epoch 8/50
8/8 [=====] - 7s 856ms/step - loss: 0.4840 - auc_3
9: 0.8529 - val_loss: 0.5143 - val_auc_39: 0.8574
Epoch 9/50
8/8 [=====] - 7s 857ms/step - loss: 0.4272 - auc_3
9: 0.8837 - val_loss: 0.4543 - val_auc_39: 0.8618
Epoch 10/50
8/8 [=====] - 7s 863ms/step - loss: 0.4018 - auc_3
9: 0.8975 - val_loss: 0.4516 - val_auc_39: 0.8612
Epoch 11/50
8/8 [=====] - 7s 879ms/step - loss: 0.4054 - auc_3
9: 0.8929 - val_loss: 0.4789 - val_auc_39: 0.8505
Epoch 12/50
8/8 [=====] - 7s 876ms/step - loss: 0.3610 - auc_3
9: 0.9310 - val_loss: 0.4669 - val_auc_39: 0.8524
Epoch 13/50
8/8 [=====] - 7s 870ms/step - loss: 0.3343 - auc_3
9: 0.9364 - val_loss: 0.5307 - val_auc_39: 0.8537
Epoch 14/50
8/8 [=====] - 7s 867ms/step - loss: 0.2894 - auc_3
9: 0.9552 - val_loss: 0.4789 - val_auc_39: 0.8501
Epoch 15/50
8/8 [=====] - 7s 876ms/step - loss: 0.2764 - auc_3
9: 0.9580 - val_loss: 0.4887 - val_auc_39: 0.8417
Training time: 112.58 seconds

```

```

In [66]: # Plot loss
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend()
plt.show()


# Calculate ROC-AUC score
y_pred_train = cnn_final.predict(train_generator)
y_pred_val = cnn_final.predict(validation_generator)
y_train = train_generator.labels
y_val = validation_generator.labels

train_roc_auc = roc_auc_score(y_train, y_pred_train.flatten())
val_roc_auc = roc_auc_score(y_val, y_pred_val.flatten())

print(f'Training ROC-AUC: {train_roc_auc:.4f}')
print(f'Validation ROC-AUC: {val_roc_auc:.4f}')

# Calculate classification matrix
y_pred_val_class = (y_pred_val > 0.5).astype('int32') # default threshold of 0.5
print(ClassificationReport(y_val, y_pred_val_class))
print(confusion_matrix(y_val, y_pred_val_class))

```



```

8/8 [=====] - 2s 243ms/step
2/2 [=====] - 1s 223ms/step
Training ROC-AUC: 0.4989
Validation ROC-AUC: 0.5126

```

	precision	recall	f1-score	support
0	0.61	0.62	0.62	151
1	0.41	0.40	0.41	99
accuracy			0.54	250
macro avg	0.51	0.51	0.51	250
weighted avg	0.53	0.54	0.54	250

```

[[ 94 57]
 [ 59 40]]

```

The tuned CNN model achieved a validation ROC-AUC score of 0.5338, demonstrating improved ability to distinguish between classes. The learning curves show a good balance between training and validation performance, with signs of potential overfitting only starting to emerge at the end of training, suggesting that the model has been effectively regularized. Overall, the tuned model's performance metrics, including accuracy, precision, and recall, indicate a strong and well-balanced performance.

Resnet CNN

Base model

```

In [59]: def residual_block(x, filters, kernel_size=(3, 3), strides=(1, 1)):
          residual = x
          if x.shape[-1] != filters:
              residual = Conv2D(filters, (1, 1), strides=strides, padding='same')(
                  x)
              x = Conv2D(filters, kernel_size, strides=strides, activation='relu', padding='same')(
                  x)
              x = Conv2D(filters, kernel_size, padding='same')(x)
              x = Add()([x, residual])
              x = Activation('relu')(x)
          return x

          def ResNet_CNN(input_shape):
              inputs = Input(shape=input_shape)
              x = Conv2D(32, (3, 3), activation='relu', padding='same')(inputs)
              x = MaxPooling2D((2, 2))(x)

              x = residual_block(x, 32)
              x = residual_block(x, 32)
              x = MaxPooling2D((2, 2))(x)

              x = residual_block(x, 64, strides=(2, 2)) # Use strides=(2, 2) to match
              x = residual_block(x, 64)
              x = MaxPooling2D((2, 2))(x)

              x = GlobalAveragePooling2D()(x)
              x = Dense(128, activation='relu')(x)
              x = Dropout(0.2)(x)
              outputs = Dense(1, activation='sigmoid')(x)

              model = Model(inputs=inputs, outputs=outputs)
              return model

```

```
In [54]: # Define constants
IMG_WIDTH, IMG_HEIGHT = 96, 96
BATCH_SIZE = 64
VALIDATION_SPLIT = 0.2

# Define data generators
train_datagen = ImageDataGenerator(rescale=1./255,
                                   validation_split=VALIDATION_SPLIT)

train_generator = train_datagen.flow_from_dataframe(
    dataframe=subsample_df,
    x_col='file_path',
    y_col='label',
    target_size=(IMG_WIDTH, IMG_HEIGHT),
    batch_size=BATCH_SIZE,
    class_mode='raw',
    subset='training')

validation_generator = train_datagen.flow_from_dataframe(
    dataframe=subsample_df,
    x_col='file_path',
    y_col='label',
    target_size=(IMG_WIDTH, IMG_HEIGHT),
    batch_size=BATCH_SIZE,
    class_mode='raw',
    subset='validation')

Found 1000 validated image filenames.
Found 250 validated image filenames.
```

```
In [ ]: resnet_model = ResNet_CNN((IMG_WIDTH, IMG_HEIGHT, 3))
resnet_model.compile(optimizer=Adam(learning_rate=0.001), loss='binary_crossentropy')
start_time = time.time()
resnet_history = resnet_model.fit(train_generator,
                                  epochs=50,
                                  validation_data=validation_generator,
                                  workers=workers,
                                  use_multiprocessing=False,
                                  callbacks=[early_stopping])
end_time = time.time()
training_time = end_time - start_time
print(f"Training time: {training_time:.2f} seconds")
```

```
Epoch 1/50
16/16 [=====] - 18s 949ms/step - loss: 0.6993 - auc_28: 0.5039 - val_loss: 0.6683 - val_auc_28: 0.6370
Epoch 2/50
16/16 [=====] - 15s 915ms/step - loss: 0.6807 - auc_28: 0.4927 - val_loss: 0.6660 - val_auc_28: 0.6649
Epoch 3/50
16/16 [=====] - 14s 868ms/step - loss: 0.6665 - auc_28: 0.5549 - val_loss: 0.6664 - val_auc_28: 0.7242
Epoch 4/50
16/16 [=====] - 14s 850ms/step - loss: 0.6653 - auc_28: 0.5573 - val_loss: 0.6592 - val_auc_28: 0.7675
Epoch 5/50
16/16 [=====] - 14s 871ms/step - loss: 0.6531 - auc_28: 0.6446 - val_loss: 0.6295 - val_auc_28: 0.8024
Epoch 6/50
16/16 [=====] - 14s 844ms/step - loss: 0.6628 - auc_28: 0.5836 - val_loss: 0.6319 - val_auc_28: 0.8217
Epoch 7/50
16/16 [=====] - 15s 874ms/step - loss: 0.6459 - auc_28: 0.6400 - val_loss: 0.6186 - val_auc_28: 0.7658
Epoch 8/50
16/16 [=====] - 14s 847ms/step - loss: 0.6532 - auc_28: 0.6124 - val_loss: 0.6717 - val_auc_28: 0.8113
Epoch 9/50
16/16 [=====] - 15s 882ms/step - loss: 0.6387 - auc_28: 0.6921 - val_loss: 0.5928 - val_auc_28: 0.8202
Epoch 10/50
16/16 [=====] - 15s 878ms/step - loss: 0.6076 - auc_28: 0.7098 - val_loss: 0.5766 - val_auc_28: 0.8551
Epoch 11/50
16/16 [=====] - 15s 874ms/step - loss: 0.6018 - auc_28: 0.7376 - val_loss: 0.5225 - val_auc_28: 0.8584
Epoch 12/50
16/16 [=====] - 14s 869ms/step - loss: 0.5334 - auc_28: 0.7961 - val_loss: 0.4789 - val_auc_28: 0.8605
Epoch 13/50
16/16 [=====] - 15s 893ms/step - loss: 0.5553 - auc_28: 0.7753 - val_loss: 0.4975 - val_auc_28: 0.8712
Epoch 14/50
16/16 [=====] - 18s 1s/step - loss: 0.5330 - auc_28: 0.7955 - val_loss: 0.4472 - val_auc_28: 0.8676
Epoch 15/50
16/16 [=====] - 24s 1s/step - loss: 0.5190 - auc_28: 0.8099 - val_loss: 0.4771 - val_auc_28: 0.8535
Epoch 16/50
16/16 [=====] - 24s 1s/step - loss: 0.5006 - auc_28: 0.8236 - val_loss: 0.4553 - val_auc_28: 0.8737
Epoch 17/50
16/16 [=====] - 17s 1s/step - loss: 0.5113 - auc_28: 0.8156 - val_loss: 0.4579 - val_auc_28: 0.8696
Epoch 18/50
16/16 [=====] - 15s 877ms/step - loss: 0.5230 - auc_28: 0.8069 - val_loss: 0.4936 - val_auc_28: 0.8729
Epoch 19/50
16/16 [=====] - 15s 880ms/step - loss: 0.5054 - auc_28: 0.8234 - val_loss: 0.4922 - val_auc_28: 0.8737
Training time: 311.28 seconds
```




```
In [ ]: # Plot loss
plt.plot(resnet_history.history['loss'], label='Training Loss')
plt.plot(resnet_history.history['val_loss'], label='Validation Loss')
plt.legend()
plt.show()

# Calculate ROC-AUC score
y_pred_train = resnet_model.predict(train_generator)
y_pred_val = resnet_model.predict(validation_generator)
y_train = train_generator.labels
y_val = validation_generator.labels

train_roc_auc = roc_auc_score(y_train, y_pred_train.flatten())
val_roc_auc = roc_auc_score(y_val, y_pred_val.flatten())

print(f'Training ROC-AUC: {train_roc_auc:.4f}')
print(f'Validation ROC-AUC: {val_roc_auc:.4f}')

# Calculate classification matrix
y_pred_val_class = (y_pred_val > 0.5).astype('int32') # default threshold of 0.5
print(classification_report(y_val, y_pred_val_class))
print(confusion_matrix(y_val, y_pred_val_class))
```



```
16/16 [=====] - 4s 218ms/step
4/4 [=====] - 1s 193ms/step
Training ROC-AUC: 0.5064
Validation ROC-AUC: 0.5051
      precision    recall  f1-score   support

     0       0.61       0.66       0.63       151
     1       0.40       0.35       0.38        99

 accuracy          0.50          0.50          0.54       250
 macro avg          0.50          0.50          0.50       250
weighted avg          0.53          0.54          0.53       250

[[99 52]
 [64 35]]
```

The ResNet CNN model achieved a validation ROC-AUC score of 0.5051, indicating that it is barely better than random chance in distinguishing between the two classes. The learning curve shows a promising trend, with both training and validation loss decreasing without signs of overfitting, suggesting that further training may improve the model's performance. However, the model's precision, recall, and F1-score metrics indicate that it struggles to accurately classify the classes, particularly class 1.

Hyper parameter tuning

- Same parameters as Simple CNN

```
In [62]: # Define hyperparameter grid
batch_sizes = [32, 64, 128]
learning_rates = [0.0001, 0.001, 0.01]

# Create a list to store results
results = []

for batch_size in batch_sizes:
    for learning_rate in learning_rates:
        # Define data generators
        train_generator = train_datagen.flow_from_dataframe(
            dataframe=subsample_df,
            x_col='file_path',
            y_col='label',
            target_size=(IMG_WIDTH, IMG_HEIGHT),
            batch_size=batch_size,
            class_mode='raw',
            subset='training')

        validation_generator = train_datagen.flow_from_dataframe(
            dataframe=subsample_df,
            x_col='file_path',
            y_col='label',
            target_size=(IMG_WIDTH, IMG_HEIGHT),
            batch_size=batch_size,
            class_mode='raw',
```

```

subset='validation')

# Define model architecture
def residual_block(x, filters, kernel_size=(3, 3), strides=(1, 1)):
    residual = x
    if x.shape[-1] != filters:
        residual = Conv2D(filters, (1, 1), strides=strides, padding='same')(x)
    x = Conv2D(filters, kernel_size, strides=strides, activation='relu')(residual)
    x = Conv2D(filters, kernel_size, padding='same')(x)
    x = Add()([x, residual])
    x = Activation('relu')(x)
    return x

def ResNet_CNN(input_shape):
    inputs = Input(shape=input_shape)
    x = Conv2D(32, (3, 3), activation='relu', padding='same')(inputs)
    x = MaxPooling2D((2, 2))(x)

    x = residual_block(x, 32)
    x = residual_block(x, 32)
    x = MaxPooling2D((2, 2))(x)

    x = residual_block(x, 64, strides=(2, 2))
    x = residual_block(x, 64)
    x = MaxPooling2D((2, 2))(x)

    x = GlobalAveragePooling2D()(x)
    x = Dense(128, activation='relu')(x)
    x = Dropout(0.2)(x)
    outputs = Dense(1, activation='sigmoid')(x)

    model = Model(inputs=inputs, outputs=outputs)
    return model

model = ResNet_CNN((IMG_WIDTH, IMG_HEIGHT, 3))

# Compile model
model.compile(optimizer=Adam(learning_rate=learning_rate), loss='binary_crossentropy')

# Define early stopping callback
early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=5,
    min_delta=0.001,
    restore_best_weights=True
)

# Train model
start_time = time.time()
history = model.fit(train_generator,
                    epochs=50,
                    validation_data=validation_generator,
                    workers=workers,
                    use_multiprocessing=False,
                    callbacks=[early_stopping])
end_time = time.time()
training_time = end_time - start_time

# Calculate ROC-AUC score
y_pred_train = model.predict(train_generator)
y_pred_val = model.predict(validation_generator)
y_train = train_generator.labels
y_val = validation_generator.labels

train_roc_auc = roc_auc_score(y_train, y_pred_train.flatten())
val_roc_auc = roc_auc_score(y_val, y_pred_val.flatten())

# Store results
row_name = f'batch{batch_size}_lr{str(learning_rate).replace(".", "")}'
results.append({
    'model_name': row_name,
    'batch_size': batch_size,
    'learning_rate': learning_rate,
    'train_time': training_time,
    'train_roc_auc': train_roc_auc,
    'val_roc_auc': val_roc_auc
})

```

Found 1000 validated image filenames.

Found 250 validated image filenames.

Epoch 1/50

32/32 [=====] - 19s 483ms/step - loss: 0.6948 - auc
_29: 0.4909 - val_loss: 0.6631 - val_auc_29: 0.7046

Epoch 2/50

32/32 [=====] - 14s 429ms/step - loss: 0.6648 - auc
_29: 0.5698 - val_loss: 0.6614 - val_auc_29: 0.7567

Epoch 3/50

32/32 [=====] - 14s 431ms/step - loss: 0.6619 - auc
_29: 0.5799 - val_loss: 0.6462 - val_auc_29: 0.7617

Epoch 4/50

32/32 [=====] - 14s 431ms/step - loss: 0.6443 - auc
_29: 0.6644 - val_loss: 0.6099 - val_auc_29: 0.7947

Epoch 5/50

32/32 [=====] - 14s 438ms/step - loss: 0.6067 - auc
_29: 0.7306 - val_loss: 0.5318 - val_auc_29: 0.8455

Epoch 6/50

32/32 [=====] - 15s 449ms/step - loss: 0.5842 - auc
_29: 0.7420 - val_loss: 0.5034 - val_auc_29: 0.8480

Epoch 7/50

32/32 [=====] - 15s 464ms/step - loss: 0.6099 - auc
_29: 0.7063 - val_loss: 0.5917 - val_auc_29: 0.8234

Epoch 8/50

32/32 [=====] - 15s 455ms/step - loss: 0.5765 - auc
_29: 0.7506 - val_loss: 0.5185 - val_auc_29: 0.8660

Epoch 9/50

32/32 [=====] - 15s 458ms/step - loss: 0.5433 - auc
_29: 0.7892 - val_loss: 0.5136 - val_auc_29: 0.8471

Epoch 10/50

32/32 [=====] - 15s 461ms/step - loss: 0.5493 - auc
_29: 0.7798 - val_loss: 0.4567 - val_auc_29: 0.8659

Epoch 11/50

32/32 [=====] - 15s 454ms/step - loss: 0.5186 - auc
_29: 0.8078 - val_loss: 0.4492 - val_auc_29: 0.8674

Epoch 12/50

```
32/32 [=====] - 14s 434ms/step - loss: 0.5307 - auc
_29: 0.7976 - val_loss: 0.4480 - val_auc_29: 0.8712
Epoch 13/50
32/32 [=====] - 15s 454ms/step - loss: 0.5374 - auc
_29: 0.7901 - val_loss: 0.4912 - val_auc_29: 0.8693
Epoch 14/50
32/32 [=====] - 15s 452ms/step - loss: 0.5345 - auc
_29: 0.7924 - val_loss: 0.4658 - val_auc_29: 0.8712
Epoch 15/50
32/32 [=====] - 15s 450ms/step - loss: 0.5174 - auc
_29: 0.8103 - val_loss: 0.4417 - val_auc_29: 0.8736
Epoch 16/50
32/32 [=====] - 15s 465ms/step - loss: 0.5231 - auc
_29: 0.8048 - val_loss: 0.4504 - val_auc_29: 0.8679
Epoch 17/50
32/32 [=====] - 15s 456ms/step - loss: 0.5127 - auc
_29: 0.8143 - val_loss: 0.4439 - val_auc_29: 0.8768
Epoch 18/50
32/32 [=====] - 15s 462ms/step - loss: 0.5297 - auc
_29: 0.7966 - val_loss: 0.4520 - val_auc_29: 0.8749
Epoch 19/50
32/32 [=====] - 15s 460ms/step - loss: 0.5101 - auc
_29: 0.8157 - val_loss: 0.4581 - val_auc_29: 0.8760
Epoch 20/50
32/32 [=====] - 15s 460ms/step - loss: 0.5212 - auc
_29: 0.8072 - val_loss: 0.4533 - val_auc_29: 0.8782
32/32 [=====] - 4s 117ms/step
8/8 [=====] - 1s 128ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
32/32 [=====] - 18s 473ms/step - loss: 0.6621 - auc
_30: 0.5795 - val_loss: 0.5961 - val_auc_30: 0.8425
Epoch 2/50
32/32 [=====] - 14s 448ms/step - loss: 0.6567 - auc
_30: 0.6269 - val_loss: 0.6629 - val_auc_30: 0.7604
Epoch 3/50
32/32 [=====] - 15s 456ms/step - loss: 0.6350 - auc
_30: 0.6574 - val_loss: 0.6179 - val_auc_30: 0.8410
Epoch 4/50
32/32 [=====] - 15s 461ms/step - loss: 0.5687 - auc
_30: 0.7679 - val_loss: 0.5463 - val_auc_30: 0.7940
Epoch 5/50
32/32 [=====] - 15s 451ms/step - loss: 0.5732 - auc
_30: 0.7603 - val_loss: 0.6297 - val_auc_30: 0.8323
Epoch 6/50
32/32 [=====] - 15s 449ms/step - loss: 0.6364 - auc
_30: 0.6908 - val_loss: 0.5682 - val_auc_30: 0.8629
Epoch 7/50
32/32 [=====] - 15s 451ms/step - loss: 0.5656 - auc
_30: 0.7801 - val_loss: 0.4649 - val_auc_30: 0.8713
Epoch 8/50
32/32 [=====] - 15s 447ms/step - loss: 0.5282 - auc
_30: 0.8030 - val_loss: 0.4628 - val_auc_30: 0.8520
Epoch 9/50
32/32 [=====] - 15s 449ms/step - loss: 0.5441 - auc
_30: 0.7920 - val_loss: 0.5308 - val_auc_30: 0.8458
Epoch 10/50
32/32 [=====] - 15s 450ms/step - loss: 0.5771 - auc
_30: 0.7522 - val_loss: 0.5403 - val_auc_30: 0.8685
Epoch 11/50
32/32 [=====] - 15s 452ms/step - loss: 0.5240 - auc
_30: 0.8036 - val_loss: 0.4913 - val_auc_30: 0.8714
Epoch 12/50
32/32 [=====] - 15s 446ms/step - loss: 0.5456 - auc
_30: 0.7904 - val_loss: 0.4705 - val_auc_30: 0.8757
Epoch 13/50
32/32 [=====] - 15s 450ms/step - loss: 0.5179 - auc
_30: 0.8107 - val_loss: 0.4429 - val_auc_30: 0.8816
Epoch 14/50
32/32 [=====] - 15s 450ms/step - loss: 0.5003 - auc
_30: 0.8249 - val_loss: 0.4688 - val_auc_30: 0.8562
Epoch 15/50
32/32 [=====] - 15s 451ms/step - loss: 0.5029 - auc
_30: 0.8217 - val_loss: 0.4391 - val_auc_30: 0.8832
Epoch 16/50
32/32 [=====] - 14s 428ms/step - loss: 0.5100 - auc
_30: 0.8172 - val_loss: 0.4503 - val_auc_30: 0.8758
Epoch 17/50
32/32 [=====] - 15s 450ms/step - loss: 0.4988 - auc
_30: 0.8270 - val_loss: 0.4578 - val_auc_30: 0.8611
Epoch 18/50
32/32 [=====] - 15s 477ms/step - loss: 0.4921 - auc
_30: 0.8305 - val_loss: 0.4900 - val_auc_30: 0.8762
Epoch 19/50
32/32 [=====] - 14s 446ms/step - loss: 0.5066 - auc
_30: 0.8189 - val_loss: 0.4268 - val_auc_30: 0.8858
Epoch 20/50
32/32 [=====] - 15s 448ms/step - loss: 0.4887 - auc
_30: 0.8355 - val_loss: 0.4827 - val_auc_30: 0.8463
Epoch 21/50
32/32 [=====] - 15s 449ms/step - loss: 0.4917 - auc
_30: 0.8327 - val_loss: 0.4475 - val_auc_30: 0.8787
Epoch 22/50
32/32 [=====] - 15s 452ms/step - loss: 0.4675 - auc
_30: 0.8505 - val_loss: 0.4091 - val_auc_30: 0.8902
Epoch 23/50
32/32 [=====] - 15s 451ms/step - loss: 0.4676 - auc
_30: 0.8514 - val_loss: 0.4348 - val_auc_30: 0.8845
Epoch 24/50
32/32 [=====] - 14s 434ms/step - loss: 0.4610 - auc
_30: 0.8554 - val_loss: 0.4068 - val_auc_30: 0.8898
Epoch 25/50
32/32 [=====] - 15s 459ms/step - loss: 0.4656 - auc
_30: 0.8527 - val_loss: 0.4262 - val_auc_30: 0.8849
Epoch 26/50
32/32 [=====] - 15s 454ms/step - loss: 0.4673 - auc
_30: 0.8508 - val_loss: 0.5610 - val_auc_30: 0.8179
Epoch 27/50
32/32 [=====] - 15s 456ms/step - loss: 0.4425 - auc
_30: 0.8689 - val_loss: 0.4188 - val_auc_30: 0.8910
Epoch 28/50
32/32 [=====] - 15s 466ms/step - loss: 0.4327 - auc
_30: 0.8766 - val_loss: 0.4220 - val_auc_30: 0.8834
```

```
Epoch 29/50
32/32 [=====] - 15s 465ms/step - loss: 0.4215 - auc
_30: 0.8821 - val_loss: 0.4162 - val_auc_30: 0.8839
32/32 [=====] - 4s 117ms/step
8/8 [=====] - 1s 120ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
32/32 [=====] - 17s 468ms/step - loss: 0.6950 - auc
_31: 0.5054 - val_loss: 0.6714 - val_auc_31: 0.5000
Epoch 2/50
32/32 [=====] - 15s 451ms/step - loss: 0.6755 - auc
_31: 0.5035 - val_loss: 0.6755 - val_auc_31: 0.5000
Epoch 3/50
32/32 [=====] - 15s 448ms/step - loss: 0.6719 - auc
_31: 0.4822 - val_loss: 0.6714 - val_auc_31: 0.5000
Epoch 4/50
32/32 [=====] - 15s 465ms/step - loss: 0.6703 - auc
_31: 0.4998 - val_loss: 0.6714 - val_auc_31: 0.5000
Epoch 5/50
32/32 [=====] - 15s 458ms/step - loss: 0.6716 - auc
_31: 0.4882 - val_loss: 0.6714 - val_auc_31: 0.5000
Epoch 6/50
32/32 [=====] - 15s 466ms/step - loss: 0.6720 - auc
_31: 0.4753 - val_loss: 0.6714 - val_auc_31: 0.5000
32/32 [=====] - 4s 120ms/step
8/8 [=====] - 1s 121ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
16/16 [=====] - 18s 975ms/step - loss: 0.6764 - auc
_32: 0.4926 - val_loss: 0.6651 - val_auc_32: 0.6840
Epoch 2/50
16/16 [=====] - 16s 959ms/step - loss: 0.6651 - auc
_32: 0.5646 - val_loss: 0.6622 - val_auc_32: 0.7179
Epoch 3/50
16/16 [=====] - 16s 945ms/step - loss: 0.6640 - auc
_32: 0.5767 - val_loss: 0.6584 - val_auc_32: 0.7387
Epoch 4/50
16/16 [=====] - 16s 983ms/step - loss: 0.6597 - auc
_32: 0.6138 - val_loss: 0.6517 - val_auc_32: 0.7706
Epoch 5/50
16/16 [=====] - 16s 946ms/step - loss: 0.6519 - auc
_32: 0.6465 - val_loss: 0.6338 - val_auc_32: 0.8146
Epoch 6/50
16/16 [=====] - 17s 1s/step - loss: 0.6354 - auc_3
2: 0.7011 - val_loss: 0.5942 - val_auc_32: 0.8500
Epoch 7/50
16/16 [=====] - 16s 971ms/step - loss: 0.5944 - auc
_32: 0.7648 - val_loss: 0.5251 - val_auc_32: 0.8641
Epoch 8/50
16/16 [=====] - 16s 978ms/step - loss: 0.5720 - auc
_32: 0.7573 - val_loss: 0.5157 - val_auc_32: 0.8635
Epoch 9/50
16/16 [=====] - 16s 982ms/step - loss: 0.5369 - auc
_32: 0.7990 - val_loss: 0.5046 - val_auc_32: 0.8687
Epoch 10/50
16/16 [=====] - 16s 964ms/step - loss: 0.5222 - auc
_32: 0.8095 - val_loss: 0.4549 - val_auc_32: 0.8703
Epoch 11/50
16/16 [=====] - 16s 934ms/step - loss: 0.5019 - auc
_32: 0.8237 - val_loss: 0.4458 - val_auc_32: 0.8674
Epoch 12/50
16/16 [=====] - 15s 934ms/step - loss: 0.5338 - auc
_32: 0.7986 - val_loss: 0.5020 - val_auc_32: 0.8697
Epoch 13/50
16/16 [=====] - 15s 936ms/step - loss: 0.5204 - auc
_32: 0.8082 - val_loss: 0.4796 - val_auc_32: 0.8687
Epoch 14/50
16/16 [=====] - 16s 981ms/step - loss: 0.5166 - auc
_32: 0.8121 - val_loss: 0.4554 - val_auc_32: 0.8651
Epoch 15/50
16/16 [=====] - 16s 944ms/step - loss: 0.5017 - auc
_32: 0.8237 - val_loss: 0.4378 - val_auc_32: 0.8734
Epoch 16/50
16/16 [=====] - 16s 974ms/step - loss: 0.5024 - auc
_32: 0.8232 - val_loss: 0.4465 - val_auc_32: 0.8715
Epoch 17/50
16/16 [=====] - 16s 953ms/step - loss: 0.4990 - auc
_32: 0.8248 - val_loss: 0.4565 - val_auc_32: 0.8639
Epoch 18/50
16/16 [=====] - 16s 970ms/step - loss: 0.5046 - auc
_32: 0.8219 - val_loss: 0.4421 - val_auc_32: 0.8709
Epoch 19/50
16/16 [=====] - 16s 935ms/step - loss: 0.5039 - auc
_32: 0.8215 - val_loss: 0.4594 - val_auc_32: 0.8755
Epoch 20/50
16/16 [=====] - 17s 997ms/step - loss: 0.5078 - auc
_32: 0.8179 - val_loss: 0.4499 - val_auc_32: 0.8698
16/16 [=====] - 4s 226ms/step
4/4 [=====] - 1s 223ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
16/16 [=====] - 18s 960ms/step - loss: 0.6721 - auc
_33: 0.5190 - val_loss: 0.6619 - val_auc_33: 0.7757
Epoch 2/50
16/16 [=====] - 20s 1s/step - loss: 0.6533 - auc_3
3: 0.6368 - val_loss: 0.6162 - val_auc_33: 0.8011
Epoch 3/50
16/16 [=====] - 16s 996ms/step - loss: 0.6163 - auc
_33: 0.7281 - val_loss: 0.5887 - val_auc_33: 0.7668
Epoch 4/50
16/16 [=====] - 16s 945ms/step - loss: 0.5925 - auc
_33: 0.7517 - val_loss: 0.5153 - val_auc_33: 0.8649
Epoch 5/50
16/16 [=====] - 16s 951ms/step - loss: 0.5512 - auc
_33: 0.7801 - val_loss: 0.5080 - val_auc_33: 0.8409
Epoch 6/50
16/16 [=====] - 17s 1s/step - loss: 0.5379 - auc_3
3: 0.7957 - val_loss: 0.4378 - val_auc_33: 0.8756
Epoch 7/50
16/16 [=====] - 16s 959ms/step - loss: 0.5386 - auc
_33: 0.7949 - val_loss: 0.4612 - val_auc_33: 0.8734
```

```
Epoch 8/50
16/16 [=====] - 16s 962ms/step - loss: 0.5237 - auc
_33: 0.8106 - val_loss: 0.4558 - val_auc_33: 0.8791
Epoch 9/50
16/16 [=====] - 15s 907ms/step - loss: 0.5165 - auc
_33: 0.8141 - val_loss: 0.5344 - val_auc_33: 0.8319
Epoch 10/50
16/16 [=====] - 16s 939ms/step - loss: 0.5431 - auc
_33: 0.7932 - val_loss: 0.4719 - val_auc_33: 0.8544
Epoch 11/50
16/16 [=====] - 16s 970ms/step - loss: 0.5181 - auc
_33: 0.8110 - val_loss: 0.4677 - val_auc_33: 0.8654
16/16 [=====] - 4s 226ms/step
4/4 [=====] - 1s 221ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
16/16 [=====] - 20s 984ms/step - loss: 4.8321 - auc
_34: 0.5162 - val_loss: 0.6656 - val_auc_34: 0.6606
Epoch 2/50
16/16 [=====] - 16s 978ms/step - loss: 0.6705 - auc
_34: 0.5384 - val_loss: 0.6726 - val_auc_34: 0.7312
Epoch 3/50
16/16 [=====] - 16s 947ms/step - loss: 0.6761 - auc
_34: 0.5174 - val_loss: 0.6549 - val_auc_34: 0.6950
Epoch 4/50
16/16 [=====] - 17s 1s/step - loss: 0.6658 - auc_3
4: 0.5513 - val_loss: 0.6514 - val_auc_34: 0.7310
Epoch 5/50
16/16 [=====] - 17s 1s/step - loss: 0.6599 - auc_3
4: 0.5931 - val_loss: 0.6339 - val_auc_34: 0.6975
Epoch 6/50
16/16 [=====] - 16s 983ms/step - loss: 0.6707 - auc
_34: 0.5450 - val_loss: 0.6599 - val_auc_34: 0.7011
Epoch 7/50
16/16 [=====] - 16s 964ms/step - loss: 0.6747 - auc
_34: 0.4747 - val_loss: 0.6711 - val_auc_34: 0.5033
Epoch 8/50
16/16 [=====] - 16s 960ms/step - loss: 0.6715 - auc
_34: 0.4903 - val_loss: 0.6710 - val_auc_34: 0.5464
Epoch 9/50
16/16 [=====] - 15s 917ms/step - loss: 0.6717 - auc
_34: 0.4716 - val_loss: 0.6691 - val_auc_34: 0.6939
Epoch 10/50
16/16 [=====] - 16s 931ms/step - loss: 0.6607 - auc
_34: 0.6077 - val_loss: 0.6617 - val_auc_34: 0.7072
16/16 [=====] - 4s 222ms/step
4/4 [=====] - 1s 234ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
8/8 [=====] - 21s 2s/step - loss: 0.7070 - auc_35:
0.5072 - val_loss: 0.6740 - val_auc_35: 0.6178
Epoch 2/50
8/8 [=====] - 19s 2s/step - loss: 0.6685 - auc_35:
0.5413 - val_loss: 0.6656 - val_auc_35: 0.6765
Epoch 3/50
8/8 [=====] - 18s 2s/step - loss: 0.6678 - auc_35:
0.5482 - val_loss: 0.6639 - val_auc_35: 0.6891
Epoch 4/50
8/8 [=====] - 18s 2s/step - loss: 0.6667 - auc_35:
0.5570 - val_loss: 0.6616 - val_auc_35: 0.6969
Epoch 5/50
8/8 [=====] - 20s 2s/step - loss: 0.6645 - auc_35:
0.5686 - val_loss: 0.6597 - val_auc_35: 0.7011
Epoch 6/50
8/8 [=====] - 19s 2s/step - loss: 0.6610 - auc_35:
0.5940 - val_loss: 0.6566 - val_auc_35: 0.7074
Epoch 7/50
8/8 [=====] - 18s 2s/step - loss: 0.6586 - auc_35:
0.6111 - val_loss: 0.6523 - val_auc_35: 0.7176
Epoch 8/50
8/8 [=====] - 18s 2s/step - loss: 0.6536 - auc_35:
0.6320 - val_loss: 0.6459 - val_auc_35: 0.7134
Epoch 9/50
8/8 [=====] - 18s 2s/step - loss: 0.6479 - auc_35:
0.6519 - val_loss: 0.6352 - val_auc_35: 0.7479
Epoch 10/50
8/8 [=====] - 18s 2s/step - loss: 0.6416 - auc_35:
0.6678 - val_loss: 0.6228 - val_auc_35: 0.7965
Epoch 11/50
8/8 [=====] - 18s 2s/step - loss: 0.6212 - auc_35:
0.7242 - val_loss: 0.5987 - val_auc_35: 0.8176
Epoch 12/50
8/8 [=====] - 18s 2s/step - loss: 0.6032 - auc_35:
0.7487 - val_loss: 0.5614 - val_auc_35: 0.8450
Epoch 13/50
8/8 [=====] - 18s 2s/step - loss: 0.5846 - auc_35:
0.7698 - val_loss: 0.5680 - val_auc_35: 0.8603
Epoch 14/50
8/8 [=====] - 18s 2s/step - loss: 0.5672 - auc_35:
0.7815 - val_loss: 0.5276 - val_auc_35: 0.8682
Epoch 15/50
8/8 [=====] - 18s 2s/step - loss: 0.5462 - auc_35:
0.8003 - val_loss: 0.4802 - val_auc_35: 0.8745
Epoch 16/50
8/8 [=====] - 18s 2s/step - loss: 0.5173 - auc_35:
0.8227 - val_loss: 0.4587 - val_auc_35: 0.8744
Epoch 17/50
8/8 [=====] - 18s 2s/step - loss: 0.5159 - auc_35:
0.8145 - val_loss: 0.4577 - val_auc_35: 0.8753
Epoch 18/50
8/8 [=====] - 18s 2s/step - loss: 0.5103 - auc_35:
0.8175 - val_loss: 0.4433 - val_auc_35: 0.8725
Epoch 19/50
8/8 [=====] - 18s 2s/step - loss: 0.5108 - auc_35:
0.8158 - val_loss: 0.4337 - val_auc_35: 0.8767
Epoch 20/50
8/8 [=====] - 18s 2s/step - loss: 0.5061 - auc_35:
0.8204 - val_loss: 0.4388 - val_auc_35: 0.8749
Epoch 21/50
8/8 [=====] - 18s 2s/step - loss: 0.5123 - auc_35:
0.8145 - val_loss: 0.4472 - val_auc_35: 0.8753
Epoch 22/50
```

```

8/8 [=====] - 18s 2s/step - loss: 0.5032 - auc_35:
0.8229 - val_loss: 0.4513 - val_auc_35: 0.8717
Epoch 23/50
8/8 [=====] - 18s 2s/step - loss: 0.5161 - auc_35:
0.8113 - val_loss: 0.4356 - val_auc_35: 0.8771
Epoch 24/50
8/8 [=====] - 18s 2s/step - loss: 0.5100 - auc_35:
0.8176 - val_loss: 0.4554 - val_auc_35: 0.8777
8/8 [=====] - 4s 481ms/step
2/2 [=====] - 1s 432ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
8/8 [=====] - 20s 2s/step - loss: 0.6741 - auc_36:
0.5304 - val_loss: 0.6595 - val_auc_36: 0.7588
Epoch 2/50
8/8 [=====] - 18s 2s/step - loss: 0.6586 - auc_36:
0.6103 - val_loss: 0.6552 - val_auc_36: 0.7739
Epoch 3/50
8/8 [=====] - 18s 2s/step - loss: 0.6568 - auc_36:
0.6130 - val_loss: 0.6199 - val_auc_36: 0.8252
Epoch 4/50
8/8 [=====] - 18s 2s/step - loss: 0.6361 - auc_36:
0.6716 - val_loss: 0.5956 - val_auc_36: 0.8462
Epoch 5/50
8/8 [=====] - 18s 2s/step - loss: 0.5886 - auc_36:
0.7818 - val_loss: 0.5153 - val_auc_36: 0.8545
Epoch 6/50
8/8 [=====] - 18s 2s/step - loss: 0.5554 - auc_36:
0.7917 - val_loss: 0.5445 - val_auc_36: 0.8020
Epoch 7/50
8/8 [=====] - 18s 2s/step - loss: 0.5486 - auc_36:
0.7839 - val_loss: 0.4925 - val_auc_36: 0.8702
Epoch 8/50
8/8 [=====] - 18s 2s/step - loss: 0.5190 - auc_36:
0.8123 - val_loss: 0.4346 - val_auc_36: 0.8776
Epoch 9/50
8/8 [=====] - 18s 2s/step - loss: 0.5254 - auc_36:
0.8109 - val_loss: 0.4826 - val_auc_36: 0.8573
Epoch 10/50
8/8 [=====] - 19s 2s/step - loss: 0.5339 - auc_36:
0.8001 - val_loss: 0.4856 - val_auc_36: 0.8522
Epoch 11/50
8/8 [=====] - 18s 2s/step - loss: 0.5199 - auc_36:
0.8081 - val_loss: 0.4521 - val_auc_36: 0.8638
Epoch 12/50
8/8 [=====] - 19s 2s/step - loss: 0.5233 - auc_36:
0.8059 - val_loss: 0.4562 - val_auc_36: 0.8762
Epoch 13/50
8/8 [=====] - 19s 2s/step - loss: 0.4917 - auc_36:
0.8300 - val_loss: 0.4520 - val_auc_36: 0.8573
8/8 [=====] - 4s 470ms/step
2/2 [=====] - 1s 429ms/step
Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
8/8 [=====] - 20s 2s/step - loss: 10.7767 - auc_37:
0.5128 - val_loss: 0.6682 - val_auc_37: 0.5821
Epoch 2/50
8/8 [=====] - 18s 2s/step - loss: 0.6672 - auc_37:
0.5396 - val_loss: 0.6730 - val_auc_37: 0.5232
Epoch 3/50
8/8 [=====] - 18s 2s/step - loss: 0.6758 - auc_37:
0.4924 - val_loss: 0.6720 - val_auc_37: 0.5000
Epoch 4/50
8/8 [=====] - 18s 2s/step - loss: 0.6720 - auc_37:
0.5176 - val_loss: 0.6720 - val_auc_37: 0.5000
Epoch 5/50
8/8 [=====] - 18s 2s/step - loss: 0.6708 - auc_37:
0.4951 - val_loss: 0.6714 - val_auc_37: 0.5000
Epoch 6/50
8/8 [=====] - 18s 2s/step - loss: 0.6711 - auc_37:
0.4964 - val_loss: 0.6714 - val_auc_37: 0.5000
8/8 [=====] - 4s 483ms/step
2/2 [=====] - 1s 424ms/step

```

```
In [63]: resnet_results = pd.DataFrame(results).set_index('model_name')
resnet_results
```

```
Out[63]:
```

	batch_size	learning_rate	train_time	train_roc_auc	val_roc_auc
--	------------	---------------	------------	---------------	-------------

model_name					
batch32_lr0001	32	0.0001	311.225434	0.495093	0.549067
batch32_lr0001	32	0.0010	442.930642	0.536388	0.491538
batch32_lr001	32	0.0100	92.291827	0.500000	0.500000
batch64_lr0001	64	0.0001	327.840774	0.516083	0.462038
batch64_lr0001	64	0.0010	186.427148	0.505087	0.455683
batch64_lr001	64	0.0100	165.473755	0.469594	0.513947
batch128_lr0001	128	0.0001	441.781146	0.495777	0.531741
batch128_lr0001	128	0.0010	240.610043	0.526558	0.460432
batch128_lr001	128	0.0100	109.543905	0.512423	0.490334

The grid search results for the ResNet model show that the best validation ROC-AUC score is 0.549067, achieved with a batch size of 32 and a learning rate of 0.0001. The training times vary significantly across different batch sizes and learning rates, with the fastest training time being 92.29 seconds for a batch size of 32 and a learning rate of 0.01. The results also show that the model's performance is sensitive to the choice of hyperparameters, with some combinations resulting in ROC-AUC scores close to 0.5. The best-performing model has a relatively low training ROC-AUC score, suggesting that it may not be overfitting. Overall, the results suggest that careful tuning of hyperparameters is necessary to achieve good performance with the ResNet model.

Comparison to Grid Search of Simple CNN

- Compared to the grid search results of the simple CNN, the ResNet model achieves a higher best validation ROC-AUC score (0.549067 vs 0.5126).
- However, the ResNet model's performance is also more sensitive to the choice of hyperparameters, with a wider range of ROC-AUC scores across different combinations.

Conclusion

Based on the results, the best model to use would be the ResNet model with a batch size of 32 and a learning rate of 0.0001. This model achieved the highest validation ROC-AUC score (0.549067) among all the models, indicating good performance on the validation set. However, it's worth noting that this model had a relatively long training time (311.23 seconds) and a low training ROC-AUC score (0.495093), which may indicate some underfitting.

If training time is a concern, the simple CNN model with a batch size of 128 and a learning rate of 0.001 might be a good alternative, as it achieved a validation ROC-AUC score of 0.529534 with a relatively shorter training time (113.99 seconds).

Given the spread and performance, the simple CNN model with a batch size of 128 and a learning rate of 0.001 seems like a reliable choice due to its relatively balanced performance and shorter training time. However, further investigation into the ResNet model's unusual performance discrepancy between training and validation might be warranted to understand if it's a one-off or a consistent pattern.

The experiment involved tuning a simple CNN and a ResNet model on a subsampled dataset due to computational constraints, testing different batch sizes and learning rates. The ResNet model achieved its best validation ROC-AUC score of 0.549067 with a batch size of 32 and a learning rate of 0.0001, while the simple CNN model achieved a best score of 0.5126 with a batch size of 128 and a learning rate of 0.001. The results show that careful tuning of hyperparameters is crucial for improving model performance. However, due to the limited scope of the grid search, which only tested batch size and learning rate, further improvements could be explored by tuning other hyperparameters such as the number of convolutional layers, dense layers, dropout rates, and activation functions. Future experiments could also benefit from testing different optimizers, regularization techniques, and data augmentation strategies. Given the resource-heavy nature of this task, exploring these additional hyperparameters and techniques could potentially lead to further performance gains.

```
In [68]: # Define constants
IMG_WIDTH, IMG_HEIGHT = 96, 96
BATCH_SIZE = 32
VALIDATION_SPLIT = 0.2

# Define data generators
train_datagen = ImageDataGenerator(rescale=1./255,
                                   validation_split=VALIDATION_SPLIT)

train_generator = train_datagen.flow_from_dataframe(
    dataframe=subsample_df,
```

```

x_col='file_path',
y_col='label',
target_size=(IMG_WIDTH, IMG_HEIGHT),
batch_size=BATCH_SIZE,
class_mode='raw',
subset='training')

validation_generator = train_datagen.flow_from_dataframe(
    dataframe=subsample_df,
    x_col='file_path',
    y_col='label',
    target_size=(IMG_WIDTH, IMG_HEIGHT),
    batch_size=BATCH_SIZE,
    class_mode='raw',
    subset='validation')

resnet_final = ResNet_CNN((IMG_WIDTH, IMG_HEIGHT, 3))
resnet_final.compile(optimizer=Adam(learning_rate=0.001), loss='binary_crossentropy')
start_time = time.time()
resnet_history = resnet_final.fit(train_generator,
    epochs=50,
    validation_data=validation_generator,
    workers=workers,
    use_multiprocessing=False,
    callbacks=[early_stopping])
end_time = time.time()
training_time = end_time - start_time
print(f"Training time: {training_time:.2f} seconds")

# Plot loss
plt.plot(resnet_history.history['loss'], label='Training Loss')
plt.plot(resnet_history.history['val_loss'], label='Validation Loss')
plt.legend()
plt.show()

# Calculate ROC-AUC score
y_pred_train = resnet_final.predict(train_generator)
y_pred_val = resnet_final.predict(validation_generator)
y_train = train_generator.labels
y_val = validation_generator.labels

train_roc_auc = roc_auc_score(y_train, y_pred_train.flatten())
val_roc_auc = roc_auc_score(y_val, y_pred_val.flatten())

print(f'Training ROC-AUC: {train_roc_auc:.4f}')
print(f'Validation ROC-AUC: {val_roc_auc:.4f}')


# Calculate classification matrix
y_pred_val_class = (y_pred_val > 0.5).astype('int32') # default threshold of 0.5
print(classification_report(y_val, y_pred_val_class))
print(confusion_matrix(y_val, y_pred_val_class))

```

```

Found 1000 validated image filenames.
Found 250 validated image filenames.
Epoch 1/50
32/32 [=====] - 17s 445ms/step - loss: 0.6798 - auc
_40: 0.5015 - val_loss: 0.6790 - val_auc_40: 0.5997
Epoch 2/50
32/32 [=====] - 14s 425ms/step - loss: 0.6668 - auc
_40: 0.5537 - val_loss: 0.6522 - val_auc_40: 0.8060
Epoch 3/50
32/32 [=====] - 13s 411ms/step - loss: 0.6451 - auc
_40: 0.6585 - val_loss: 0.5823 - val_auc_40: 0.8262
Epoch 4/50
32/32 [=====] - 14s 417ms/step - loss: 0.6134 - auc
_40: 0.7043 - val_loss: 0.6362 - val_auc_40: 0.6459
Epoch 5/50
32/32 [=====] - 14s 434ms/step - loss: 0.6441 - auc
_40: 0.5998 - val_loss: 0.6021 - val_auc_40: 0.7335
Epoch 6/50
32/32 [=====] - 14s 426ms/step - loss: 0.6082 - auc
_40: 0.7123 - val_loss: 0.5548 - val_auc_40: 0.8597
Epoch 7/50
32/32 [=====] - 14s 421ms/step - loss: 0.5772 - auc
_40: 0.7505 - val_loss: 0.5926 - val_auc_40: 0.8279
Epoch 8/50
32/32 [=====] - 14s 431ms/step - loss: 0.6029 - auc
_40: 0.7405 - val_loss: 0.5030 - val_auc_40: 0.8641
Epoch 9/50
32/32 [=====] - 14s 432ms/step - loss: 0.5289 - auc
_40: 0.8033 - val_loss: 0.4937 - val_auc_40: 0.8721
Epoch 10/50
32/32 [=====] - 14s 434ms/step - loss: 0.5291 - auc
_40: 0.7975 - val_loss: 0.4682 - val_auc_40: 0.8737
Epoch 11/50
32/32 [=====] - 15s 440ms/step - loss: 0.5297 - auc
_40: 0.8001 - val_loss: 0.4652 - val_auc_40: 0.8747
Epoch 12/50
32/32 [=====] - 15s 466ms/step - loss: 0.5044 - auc
_40: 0.8198 - val_loss: 0.4286 - val_auc_40: 0.8800
Epoch 13/50
32/32 [=====] - 15s 462ms/step - loss: 0.4968 - auc
_40: 0.8273 - val_loss: 0.4411 - val_auc_40: 0.8734
Epoch 14/50
32/32 [=====] - 15s 451ms/step - loss: 0.4935 - auc
_40: 0.8294 - val_loss: 0.4378 - val_auc_40: 0.8744
Epoch 15/50
32/32 [=====] - 15s 445ms/step - loss: 0.4802 - auc
_40: 0.8416 - val_loss: 0.4861 - val_auc_40: 0.8595
Epoch 16/50
32/32 [=====] - 15s 446ms/step - loss: 0.5095 - auc
_40: 0.8159 - val_loss: 0.4748 - val_auc_40: 0.8688
Epoch 17/50
32/32 [=====] - 15s 446ms/step - loss: 0.5111 - auc
_40: 0.8166 - val_loss: 0.4561 - val_auc_40: 0.8690
Training time: 253.25 seconds

```

```
32/32 [=====] - 4s 113ms/step
8/8 [=====] - 1s 121ms/step
Training ROC-AUC: 0.4815
Validation ROC-AUC: 0.5445
```

	precision	recall	f1-score	support
0	0.62	0.66	0.64	151
1	0.42	0.37	0.40	99
accuracy			0.55	250
macro avg	0.52	0.52	0.52	250
weighted avg	0.54	0.55	0.54	250

```
[[100  51]
 [ 62  37]]
```

```
In [67]: test_datagen = ImageDataGenerator(rescale=1./255)
```

```
test_generator = test_datagen.flow_from_dataframe(
    dataframe=test_df,
    x_col='file_path',
    target_size=(IMG_WIDTH, IMG_HEIGHT),
    batch_size=BATCH_SIZE,
    class_mode=None,
    shuffle=False)
```

Found 57458 validated image filenames.

```
In [71]: test_predictions = resnet_final.predict(test_generator)
test_binary_predictions = (test_predictions > 0.5).astype(int)
submission_df = pd.DataFrame({
    'id': [path.split('/')[-1].split('.')[0] for path in test_generator.file
    'label': test_binary_predictions.flatten()
})
```

```
449/449 [=====] - 209s 465ms/step
```

```
In [72]: submission_df
```

```
Out[72]:
```

	id	label
0	fd0a060ef9c30c9a83f6b4bfb568db74b099154d	1
1	1f9ee06f06d329eb7902a2e03ab3835dd0484581	1
2	19709bec800f372d0b1d085da6933dd3ef108846	1
3	7a34fc34523063f13f0617f7518a0330f6187bd3	0
4	93be720ca2b95fe2126cf2e1ed752bd759e9b0ed	0
...
57453	2581931c6ef068f105a872f2c5500275fc678242	0
57454	11b250a664d09ab59fd2afbdb2f8d786763b185d	1
57455	18a6030935ec1ef1ce486ec51bc95abb4008fbf1	0
57456	f541404e501e23a0188c852eb37eac94053cfdc0	0
57457	3cb6f5e2db8ad046c946b581fa12d20df5ce2927	0

57458 rows x 2 columns

```
In [73]: submission_df.to_csv('submission.csv', index=False)
```

```
In [ ]:
```